

## Concurrency Issues

- All aspects of scheduling apply to tasks, threads, and interrupts
  - We preempt execution
- Must maintain deadline required in the face of all conflicts
  - The conflicts are caused by:
    - Aperiodic events disrupt execution flow
    - Deadlock & infinite loops
      - Managed (in hardware) with a watchdog timer
      - On overflow, reset the microcontroller
      - Code must reset timer to avoid this

## Resource Contention:

- 1 task/interrupt locks a resources needed by another to complete
  - Meet deadline
- Locks:
  - disable all interrupt until done (i.e. SIGNAL)
  - CPU priority upped to limit which interrupts can run(i.e. ISR)
- If we allow higher priority tasks to override lower priority ones, we can have a potential problem:
  - If a lower priority task has a lock on a resource we need, we are stuck
    - Priority inversion(name/technical term of this problem)

## Solutions:

- a) Priority inheritance
  - Lower priority task gets a higher priority temporary so it can finish
    - Returns to lower priority when it's done
  - Computationally expensive
    - OS constantly working to detect
- b) Priority ceiling
  - Define the priority required for a task to ensure that it's non-blocking
  - task/interrupt automatically gets this priority
    - When it locks the resource
  - Priority based on highest priority task that can lock it

**Priority inversion could occur in A3, need to analyze and implement appropriate priority ceiling**

---

How can we transfer data out of an interrupt without locking?

- Read and write memory without priority problems

4-slot algorithm:

- Want concurrency read/write of consistent data
  - Not necessarily "latest"

- Maintain 2 pairs of data
  - Each pair has 2 slots so we can write to 1 while reading the other
  - IMG 81
  - Read:
    - Read from the lastWritten pair
      - After read we set lastRead to point to this pair
    - Use currentSlot for that pair
  - Write:
    - Write to the opposite pair from lastRead
    - Use the opposite slot for that pair
    - After write, update currentSlot and lastWritten to where we just wrote
  - If we have concurrent read & write, we will access the same pair but different slots

### Scheduling

- We have a number of variables to consider:
  - Preemptive vs non-preemptive system
  - Static vs dynamic scheduling
  - Aperiodic tasks
  - Load on CPU
- Preemption
  - Do we allow 1 task to stop the execution of another?
    - How does that affect the deadline?
      - Should we preempt a task with a hard deadline?
  - Note: non-preemptive scheduling is deterministic
- Static vs dynamic:
  - Can tasks be introduced on the fly?
    - Do we fix the task and their execution order?
      - Most common solution
  - Note: we can have static with preemption
- Aperiodic tasks:
  - Having just periodic tasks is “much” easier
  - When an aperiodic task arrives, how do we handle it?
    1. Insert as another task with regular scheduling
      - Exit after 1 run
    - Note: a scheduler is just iterating over a table of function pointers
      - Then, aperiodic task is just another call
    2. Allow aperiodic task to run when no periodic task need to run.
      - Known as slack stealing
    3. Make them solely interrupt driven
      - Good for hard deadlines

### Loading:

- As tasks are added (or take more CPU time), how do we respond?

- Must prioritize hard deadlines (soft deadlines could be missed compare to hard deadlines)
  - As load increases, preemption becomes important (when multiple hard deadline competing)
    - To the point of using fail safe or “limp home” mode
- We design for load ahead of time
  - Instead of a scheduler with preemption, build a schedule that has load guarantees